Screen View Protocol

Josh Brown

December 11, 2021

1 Abstract

Screen View is an end to end encrypted remote screen viewing and controlling software. This document describes the protocols necessary to make it function.

Contents

2　Definitions

The following definitions are used globally throughout the document:
The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in RFC 2119.

3 Introduction

This section describes the high level protocol for Screen Viewincorporating all the protocols defined below. Terms used in this section are defined in other sections.

3.1 Application Layers

The table belows is an abstract diagram of the different layers of Screen View. Each layer encapsulates all the below layers.

Transport Layer
| Server Encryption Layer
|| Server Communication Layer
||| E2EE (Peer $\leftrightarrow$ Peer) Encryption Layer
|||| Host $\leftrightarrow$ Client Communication Layer

The Transport Layer is the OSI Transport Layer level protocol for networking (TCP or UDP).

The Server Encryption layer provides security between the Server and a Peer (Client and Host). The Server Encryption Layer is discussed in 4.

3.2 Initialization

To begin, a Peer connects to the server over TCP and commences the Server Encryption Layer protocol for TCP defined in 4. Next the Server Communication Layer defined in section 5 commences. All Server Communication Layer messages are encrypted and encapsulated in messages defined in the Server Encryption Layer. Once a session is established between two peers, the E2EE (Peer $\leftrightarrow$ Peer) Encryption Layer commences as defined in 6. All E2EE (Peer $\leftrightarrow$ Peer) Encryption messages are encapsulated in Server Communication Layer messages. Finally, the Host $\leftrightarrow$ Client Communication Layer commences as defined in 6. All Host $\leftrightarrow$ Client Communication Layer messages are encrypted and encapsulated in E2EE (Peer $\leftrightarrow$ Peer) Encryption Layer messages.

3.3 Transport Messages Total Header Sizes

4 Server Encryption Layer

The Server Encryption Layer provides security for communication between Peers and the Server. TCP and UDP have different security methods.

4.1 TCP

The TCP Server Encryption Layer is heavily based on a simplification of TLS 1.3 as defined in.

4.1.1 PeerHello

Peer → Server

Bytes	Name	Value
1	type	1
16	public-key	

$(E_{peer}^{pub}, E_{peer}^{priv}) :=$ DH-Generate()

public-key $:= E_{peer}^{pub}$

4.1.2 ServerHello

Server → Peer

Bytes	Name	Value
1	type	2
3	certificates-length	
certificates-length	certificate_list	
16	public-key	
variable	certificate-verify	

certificate_list is defined in RFC8446 Section-4.4.2.

$$(E_{serv}^{pub}, E_{serv}^{priv}) := \text{DH-Generate()}$$
$$\text{public-key} := E_{serv}^{pub}$$

certificate-verify is defined in RFC8445 Section-4.4.3 with the following modification. The content that is signed is:

$$\text{content} := \text{"SreenViewServerVerify"} \mathbin{||} 0 \mathbin{||} E_{serv}^{pub}$$

The Client MUST validate all signatures in accordance with the TLS spec.

4.1.3 Transport Data Key Derivation

$$C_{peer} = \text{DH}(E^{pub}_{serv}, E^{priv}_{peer})$$

$$C_{serv} = \text{DH}(E^{pub}_{peer}, E^{priv}_{serv})$$

$$(T^{send}_{peer} = T^{recv}_{serv}, T^{recv}_{peer} = T^{send}_{serv}) := \text{KDF}(C_{peer} = C_{serv}, \epsilon)$$

$$N^{send}_{peer} = N^{recv}_{serv} = N^{recv}_{peer} = N^{send}_{serv} := 0$$

4.1.4 Subsequent Messages: Transport Data Messages

Peer $\leftrightarrow$ Server

Bytes	Name	Value
1	type	3
2	data-length	
data-length	data	

$$\text{data} := \text{AEAD}(T^{send}_m, N^{send}_m, P, \epsilon)$$

$$N^{send}_m := N^{send}_m + 1$$

Where P is the payload to be transported

4.2 UDP

UDP encryption and authentication rely on the *session-id*, *peer-id* and *peer-key* values established in a session (described in 5.4). The Server MUST NOT process or reply to any messages that don't pass authentication. This prevents an amplification attack.

4.2.1 Transport Data Key Derivation

$$G := \text{session-i}$$

$$H := \text{peer-id}$$

$$J := \text{peer-key}$$

$$(V^{send}_{peer} = T^{recv}_{serv}, P^{recv}_{peer} = P^{send}_{serv}) := \text{KDF}(\text{HASH}(G \,\|\, H \,\|\, J), \epsilon)$$

$$M^{send}_{peer} = M^{recv}_{serv} = M^{recv}_{peer} = M^{send}_{serv} := 0$$

4.2.2 Transport Data Messages

Peer $\rightarrow$ Server

Bytes	Name	Value
1	type	4
16	*peer-id*	
8	counter	
UDP length $-$ 8 bytes	data	

Server $\rightarrow$ Peer

Bytes	Name	Value
1	type	5
8	counter	
UDP length $-$ 8 bytes	data	

$$\text{data} := \text{AEAD}(V_m^{send}, M_m^{send}, P, \epsilon)$$

$$\text{counter} := M_m^{send} \; M_m^{send} := M_m^{send} + 1$$

Where P is the payload to be transported

5 Screen View Server Communication (SVSC) Protocol

The SVSC protocol is the Server Communication Layer protocol used for Peers to interact with the relay server, Server. Peers can lease an ID as well as begin a session with another Peer. Once a session is established, Peers can forward messages to another Peer. Unless otherwise noted, all messages MUST occur over TCP.

With the exception of the *Handshake* messages. All SVSC messages' first byte contain a number to indicate the message type.

5.1 Definitions

The following definitions are used globally throughout the document:

- Peer - denotes a client in classical server/client environment

- Server - The intermediary server used for routing and proxying data between

5.2 Handshake

5.2.1 ProtocolVersion

Handshaking begins by the Server sending the Peer a *ProtocolVersion* message. This lets the server know the version supported by the Host.

The *ProtocolVersion* message consists of 12 bytes interpreted as a string of ASCII characters in the format "SVSC xxx.yyy" where xxx and yyy are the major and minor version numbers, padded with zeros.

<div align="center">

Server → Peer

Bytes	Name	Value
11	version	"SVSC 001.000"

</div>

The Peer replies back either `0` to indicate the version is not acceptable and that the handshake has failed or `1` if the version is acceptable to the Peer and the handshake as succeeded. If 0 is sent, all communication MUST cease and the TCP connection MUST be terminated.

<div align="center">

Peer → Server

Bytes	Name	Value
1	ok	0 or 1

</div>

5.3 Leasing

A lease is a temporary assignment of an ID to a Peer. The ID format and generation is discussed in 4.2.5. A maximum of 1 ID can be leased per TCP connection. ID generation MUST be rate limited to prevent ID exhaustion. Rate limiting rules are out of scope for this protocol, however some suggestions are listed in 4.2.6.

5.3.1 LeaseRequest

A *LeaseRequest* message requests a lease of an ID.

<div align="center">

Peer → Server

Bytes	Name	Value
1	type	1
1	has-cookie	0 or 1
Below only if *has-cookie* is 1		
24	cookie	

</div>

If a Peer would like to request an ID it had previously been issued after expiration, it may include the cookie value it received in the *LeaseResponse*.

5.3.2 LeaseResponse

A *LeaseResponse* message is a response to a *LeaseRequest*.
If *has-cookie* is 1, a Server MAY consider the *cookie* value in *LeaseRequest* or completely ignore it.

<div align="center">

Server → Peer

Bytes	Name	Value
1	type	2
1	accepted	0 or 1
Below only if *accepted* is 1		
4	id	
24	cookie	
8	expiration	

</div>

expiration is a 64 bit Unix timestamp representing the expiry of lease. Disconnection of a Peer (e.g, the TCP connection is dropped) does not end a lease.

cookie a 128 bit value. The generation of this value is discussed in 4.2 .7.

Consideration of the *cookie* value MUST have no effect on the the value of *accepted*. That is, if the request is for a specific ID (implied by the presence of a cookie value and a *has-cookie* value equal to 1 in the *LeaseRequest*) and the ID requested is not available, the Server SHOULD respond with a different available ID and an *accepted* value of 1 (assuming an ID is available).

accepted SHOULD only be 0 if no IDs are left or for rate limiting reasons.

5.3.3 LeaseExtensionRequest

A *LeaseExtensionRequest* message is used to extend a lease. Before a lease has expired, the Peer can request a lease extension. The server can accept or deny this request.

Peer → Server

Bytes	Name	Value
1	type	3
24	cookie	

5.3.4 LeaseExtensionResponse

A *LeaseExtensionResponse* message is a response to a *LeaseExtensionRequest*.

Server → Peer

Bytes	Name	Value
1	type	4
1	extended	0 or 1
Below only if *extended* is 1		
8	new-expiration	

new-expiration is a 64 bit Unix timestamp representing the expiry of lease.

5.3.5 ID Generation

An ID is a 26 to 33 bit decimal number. This comes out to about up to 8 to 10 decimal digits, respectively. The Server may scale the keyspace depending on current usage. For optimal user experience while maintaining efficiency, the Server MUST only use keyspaces between 26 bits and 33 bits for ID generation. ID generation must also be uniformly random. All active IDs must be stored on the server. ID generation MAY occur using the below algorithm:

Let S represents a set of all active IDs, B be a number of bits between 26 and 33, and $generate(x)$ be a functions that returns a x uniformly random bits.

Algorithm 1 ID generation
---
   id
   repeat
      $id \leftarrow generate(B)$
   until $id \notin S$
   $S \leftarrow S \cup \{id\}$
   return id

5.3.6 Rate Limits

To prevent ID exhaustion, rate limits SHOULD be in place. TCP is used for *LeaseRequests* so IP addresses can not be spoofed. However, using proxy services such as Tor, simple IP based rate limits are likely not entirely sufficient. Servers MAY want to block all known proxy IP addresses. Additionally, a Server MAY only want to allow one active ID per Peer.

5.3.7 Cookie Value

A *cookie* value is a 128 bit value used for authentication in *LeaseExtensionRequest* and *LeaseRequest* messages. Specific generation of a *cookie* is out of scope, however care must be taken to ensure it is not predictable or exploitable. This value MAY be simply a random 24 byte key, HMAC-SHA1(id, key) || id, or something else entirely.

5.4 Sessions

A session is a connection between two Peers. At least one Peer must have an ID. A Peer can have a maximum of one session at any time. Immediately after receiving a *EstablishSessionResponse* message with a *status* of 1 or a *EstablishSessionNotification* message a Peer MUST begin the UDP Heartbeat procedure defined in TODO and continue for the duration of the session. Failure to do so MAY result in dropped *SessionData** packets.

5.4.1 EstablishSessionRequest

An *EstablishSessionRequest* message is a Peer request to establish a connection to a Peer.

Peer → Server

Bytes	Name	Value	Description
1	type	5	
4	lease-id		The ID of the Peer to establish this connection with

5.4.2 EstablishSessionResponse

An *EstablishSessionResponse* message is a response to *EstablishSessionRequest*.

Server → Peer

Bytes	Name	Value	Description
1	type	6	
4	lease-id		the ID of the Peer attempted to connect to
1	status	0-3	described below
Below only if *status* is 0			
16	session-id		described below
16	peer-id		described below
16	peer-key		described below

13

status can have the following values:

Value	Description
0	session establishment was successful
1	ID not found
2	Peer is offline
3	Peer is busy
4	You are busy
3	Other error

A Peer may be considered offline if, for example, an unexpired ID has been assigned to them and then the TCP connection is dropped. Heartbeat messages and timeouts are discussed in .

session-id is a 128 bit random value used for session identification

peer-id is a 128 bit random value used to authentication a peer for a given session.

5.4.3 EstablishSessionNotification

A *EstablishSessionNotification* notifies a Peer that a session has been established with them.

Server → Host

Bytes	Name	Value	Description
1	type	7	
16	session-id		described in 4.3.2
16	peer-id		described in 4.3.2
16	peer-key		described in 4.3.2

5.4.4 SessionEnd

A *SessionEnd* message is used to terminate a session. Once a Server receives a *SessionEnd* message,

Peer → Server

Bytes	Name	Value	Description
1	type	8	

5.4.5 SessionEndNotification

A *SessionEndNotification* notifies a Peer that a session has ended. If a Client sends a *SessionEnd* message, the Server MUST send a *SessionEndNotification* message to a Host.

Server → Peer			
Bytes	**Name**	**Value**	**Description**
1	type	9	

5.4.6 SessionDataSend - TCP/UDP

A *SessionDataSend* is a message from a Peer intended to be forwarded to the Peer on the other side of the session.

Peer → Server			
Bytes	**Name**	**Value**	**Description**
1	type	10	
3	data-length		length of the content in bytes
data-length	data		data to be forwarded

5.4.7 SessionDataReceive - TCP/UDP

A *SessionDataReceive* is a message being forwarded to a Peer from the Peer on the other side of the session. The Server SHOULD forward the message along the same transport as it was received.

Server → Peer			
Bytes	**Name**	**Value**	**Description**
1	type	11	
3	data-length		length of the content in bytes
data-length	data		data to be forwarded

6 Remote Visual Display (RVD) Protocol

The RVD protocol is used to communicate mouse input, keyboard input, frame data, and clipboard data between the Host and the Client.

All messages can occur over either TCP or UDP but it is strongly RECOMMENDED that the noted transport protocol is used.

With the exception of the *Handshake* messages. All RVD messages' first byte contain a number to indicate the message type.

6.1 Definitions

- Host - A peer with an ID that wants to share their screen to the Client

- Client - A peer that wants to view and maybe control the Host's screen

- Display - A rectangular visual region that is shared by a Host to a Client. May or may not be

- *Controllable.*

- Controllable - A *Display* that accepts keyboard and mouse input.

6.2 Handshake

6.2.1 ProtocolVersion - TCP

Handshaking begins by the Host sending the client a *ProtocolVersion* message. This lets the Client know the verison supported by the Host.

The *ProtocolVersion* message consists of 11 bytes interpreted as a string of ASCII characters in the format "RVD xxx.yyy" where xxx and yyy are the major and minor version numbers, padded with zeros.

Host → Client

Bytes	Name	Value
11	version	"RVD 001.000"

The client replies back either `0` to indicate the version is not acceptable and that the handshake has failed or `1` if the version is acceptable to the client and the handshake as succeeded. If `0` is sent, all communication MUST cease and an error SHOULD be displayed to user.

Client → Host

Bytes	Name	Value
1	ok	0 or 1

6.2.2 Initialization

Once the handshake has succeeded the Host responds with a *DisplayChange* message.

6.3 Control messages

Control messages are messages that instruct client about changes regarding the Host.

6.3.1 DisplayChange - TCP

A *DisplayChange* message informs the client about the available *Display*s. RVD supports up to 255 displays.

Host → Client

Bytes	Name	Value
1	type	1
1	clipboard-readable	0 or 1
1	number-of-displays	1 - 255
variable	displays-information	*DisplayInformation*

Each *Display* has an associated *DisplayInformation*. *displays-information* contains *number-of-displays DisplayInformation*'s. A *DisplayInformation* is defined below:

Bytes	Name	Description
1	display-id	
2	width	number of pixels of the width of this display
2	height	number of pixels of the width of this display
2	cell-width	number of pixels of the height of a cell in the grid
2	cell-height	number of pixels of the height of a cell in the grid
1	access	defined below
1	name-length	length of the *display-name* in bytes
name-length	display-name	the display name (UTF-8)

Restrictions:

- *cell-width* MUST be less than *width*. *cell-height* must be less than *height*.

- The *access* byte defines what type of access is available for the display. The bits of the

- *access* byte are described below in little endian.

17

Bit	Name
0	Flush
1	*Controllable*
2	
3	
4	Reserved for future
5	use
6	
7	

If the *Controllable* bit is 1 and the *clipboard-readable* byte is set to 1, then the clipboard is writable. The *Controllable* bit SHOULD be consistent throughout all displays.

The *Flush* bit indicates whether this display has changed, specifically if this *display-id* refers to a different *Display* than the same *display-id* did in the previous *DisplayChange* message. In initialization, this MUST always be 1 (as there is no previous *DisplayChange*). If the display hasn't changed (0) then the frame data may be maintained. If *Flush* is 0, *width*, *height* MUST remain the same as the previous *DisplayChange* specified for the *display-id*.

6.3.2 DisplayChangeReceived - TCP

The *DisplayChangeReceived* message is sent in reply after receiving a *Display-Change* message. It indicates to the Host they may start sending *FrameData* referencing the new *DisplayInformation* in the most recent *DisplayChange*.

Client → Host

Bytes	Name	Value
1	type	2

6.3.3 MouseLocation - TCP/UDP

The *MouseLocation* message send information about where the mouse is currently on the screen. The Host sends this information periodically throughout the session. The Host SHOULD send a *MouseLocation* update when mouse input is received from the Host's system or in reply when it receives a *MouseInput*.

Host → Client

Bytes	Name	Value	Description
3	type	3	
1	display-id	0-255	
2	x-location		x coordinate of the mouse
2	y-location		y coordinate of the mouse

6.4 Input

Input messages (including *MouseLocation*) may be sent over TCP or UDP. TCP is preferred in most situations. However, in situations where speed is prioritized over the guarantees TCP provides (such as gaming), UDP can be used.

6.4.1 MouseInput - TCP/UDP

Client → Host

Bytes	Name	Value	Description
1	type	4	
1	display-id	0-255	
2	x-position		x coordinate of the mouse
2	y-position		y coordinate of the mouse
1	button-mask		described below

Indicates either pointer movement or a pointer button press or release. The pointer is now at (x-position, y-position), and the current state of buttons 1 to 8 are represented by bits 0 to 7 of button-mask respectively, 0 meaning up, 1 meaning down (pressed).

On a conventional mouse, buttons 1, 2 and 3 correspond to the left, middle and right buttons on the mouse. On a wheel mouse, each step of the wheel is represented by a press and release of a certain button. Button 4 means up, button 5 means down, button 6 means left and button 7 means right.

6.4.2 KeyInput - TCP/UDP

The *KeyInput* event sends key presses or releases.

Client → Host

Bytes	Name	Value	Description
1	type	5	
1	down-flag	0 or 1	indicates whether the key is now pressed or released
4	key		"keysym"

Details can be found at the RFB Spec

6.5 Clipboard

6.5.1 ClipboardTypeRequest - TCP

Used to request clipboard types the Host supports.

Client → Host

Bytes	Name	Value
1	type	6

6.5.2 ClipboardTypeResponse - TCP

Response to the *ClipboardTypeRequest*

Host → Client

Bytes	Name	Value	Description
1	type	7	
1	number-of-clipboard-types	0-255	
variable	data (variable bytes)		described below

number-of-clipboard-types is always 0 if *clipboard-readable* is 0.

data contains *number-of-clipboard-types ClipboardType*s. *ClipboardType* is defined below.

Bytes	Name	Value	Description
1	type-length	1-255	
type-length	type-name		type name in ASCII

6.5.3 CopyRequest - TCP

This is a request for a keyboard contents. It can be made by either the Client or the Host.

Client ↔ Host

Bytes	Name	Value	Description
1	type	8	
1	type-length	1-255	
type-length	type-name		type name in ASCII

6.5.4 CopyResponse - TCP

CopyResponse message is a response to a *CopyRequest.*

Client ↔ Host

Bytes	Name	Value	Description
1	type	9	
1	accepted	0 or 1	
Below only if *accepted* is 1			
1	type-length	1-255	
type-length	type-name		type name in ASCII
3	content-length		the length of the content (maximum 2^{24} bytes or 16MB)
content-length	data		zlib'ed raw data

accepted indicates whether the *CopyRequest* was accepted. If 0, the rest of the message MUST not exist. If *clipboard-readable* is 0, *accepted* is always 0. A Client or Host may send this message without a request. If a *CopyResponse* is unsolicited, then *accepted* MUST be 1.

data is zlib compressed.

6.5.5   A note on Pasting

There is a no paste message. To paste data an unsolicited *CopyResponse* may be sent and then the keyboard shortcut (ctrl+v or cmd+v) should be sent via the *KeyboardMessage*

6.6   FrameData - UDP

The *FrameData* message contains an update of a particular cell on a particular *Display*.

Bytes	Name	Value
1	type	10
4	sequence-number	
1	display-id	0-255
2	cell-number	
2	size	
size	data	

sequence-number is an incrementing 32-bit counter for each *FrameData* sent
data contains jpeg pixel data of the updated cell.

7 Weak Pre Shared Key, Key Authentication (WPSKKA) Protocol

7.1 Introduction

This protocol is used to establish end to end encryption between the Host and the Client.

In recent years end-to-end encryption has risen in popularity due to privacy and security concerns. However, many implementations of end-to-end encryption rely on a third party and/or are susceptible to man-in-the-middle attacks making them inadequate.

Screen sharing applications are used for a multitude of different purposes. One common use case is an IT professional assisting somebody by remotely viewing and controlling their computer. Sensitive data could be visible on the user's screen. Therefore, end-to-end encryption is preferred.

Additionally, IT professionals will often be communicating with the user via the telephone. This provides a bi-directional external channel to transfer information. However, people cannot and will not transfer large amounts or complicated data reliably. Simply communicating letters can be confusing. "B", "C", "D", "E", "G" all sound similar and can be confused. Therefore, communicating strictly numbers is ideal.

This creates an issue. Short, purely numeric keys provide extremely low entropy. A 10 digit numerical code only provides about 33-bits of entropy. In fact, in order to get the ideal 128-bits of entropy a 39 digit key would need to be used. Users will not want to relay 39 digits over the phone.

One common solution when using a weak password is using a KDF to perform key-stretching. However, a 10 digit numerical code has such a small key space that it is relatively easy to brute force. Government agencies, such as the NSA, or large companies, such as Microsoft or Google could easily brute-force even very slow KDFs such as argon2 or PBKDF2.

7.2 Existing Practice

Based on TeamViewer's security statement, TeamViewer is end-to-end encrypted. However, a rogue or malicious TeamViewer intermediary server could easily provide a different set of keys to each party and intercept all communication. The parties must trust this third party in order to achieve proper end-to-end encryption. TeamViewer also does not give the ability for users to check the fingerprint of the other party's public key.

Other applications such as Zoom and Signal (and the Signal protocol itself) do give users the ability to verify some sort of fingerprint of the public key (Zoom and Signal). However, most users don't bother confirming the numbers. Again trusting the third party server.

Solutions such as TLS do provide strong end-to-end encryption but rely on a third party Certificate Authority to sign public keys. This would not be possible in the aforementioned use case.

7.3 Goal

Elliptic Curve Diffie-Hellman provides a sufficiently secure means of arriving at a shared secret. The goal of this protocol is to authenticate Elliptic Curve public keys using a weak (3 bytes length, 24 bits of entropy) pre-shared key communicated via an external channel. The security of the external channel is out of scope for this protocol and will be assumed to be secure (see Security Considerations section).

7.4 Requirements

Communication between the parties should be minimized.

The protocol should be secure in the case of a malicious Client and/or malicious Server. This means the protocol should not be susceptible to man-in-the-middle (MITM) attacks, and the Host should be able to self authenticate the Client without trusting the Server.

7.5 Other Protocol/Algorithms Definitions

The Secure Remote Password (SRP) protocol is defined in RFC2945 and RFC5054.

Elliptic Curve Diffie-Hellman (ECDH) key exchange is described in RFC6090.

7.6 Protocol

This protocol occurs after a connection is established between the Host and Client using the Server.

The SRP group is the 2048-bit group from RFC5054:

The hexadecimal value for the prime is:

```
AC6BDB41 324A9A9B F166DE5E 1389582F AF72B665 1987EE07 FC319294
3DB56050 A37329CB B4A099ED 8193E075 7767A13D D52312AB 4B03310D
CD7F48A9 DA04FD50 E8083969 EDB767B0 CF609517 9A163AB3 661A05FB
D5FAAAE8 2918A996 2F0B93B8 55F97993 EC975EEA A80D740A DBF4FF74
7359D041 D5C33EA7 1D281E44 6B14773B CA97B43A 23FB8016 76BD207A
436C6481 F1D2B907 8717461A 5B9D32E6 88F87748 544523B5 24B0D57D
5EA77A27 75D2ECFA 032CFBDB F52FB378 61602790 04E57AE6 AF874E73
03CE5329 9CCC041C 7BC308D8 2A5698F3 A8D0C382 71AE35F8 E9DBFBB6
94B5C803 D89F7AE4 35DE236D 525F5475 9B65E372 FCD68EF2 0FA7111F
9E4AFF73
```

The generator is: 2.

The Host generates the following:

- PK_H/pk_H- Host ephemeral elliptic curve Public/Private key using the `secp521r1` (P-521) curve

- I - 128 bit cryptographical secure random number, used as the identity or username in SRP

- S - SRP salt

- P - 3 byte random cryptographical secure random number, used as the password in SRP

- V - SRP verifier

- b - SRP random private value

- B - SRP public value

- k - SRP K value

The Host sends S, I, and B to the Client.

The Client generates:

- PK_C/pk_C- Client ephemeral elliptic curve Public/Private key using the `secp521r1` (P-521) curve

- a - SRP random private value

- A - SRP public value

- u - SRP u value

- k - SRP k value

- x - SRP x value (hashed P value communicated externally to the client)

- L - SRP session key

The Client sends the Host A, PK_C and HMAC(PK_C, L).

The Host derives:

- u - SRP U value

- L - SRP session key

The Host authenticates PK_C HMAC using L. If authentication is successful, the Host performs DHKE to derive, T the shared secret.

The Host sends the client PK_H and HMAC(PK_H, L).

The Client authenticates PK_H HMAC using L. If authentication is successful, the Client performs DHKE to derive, T the shared secret.

Both the Client and the Host encrypt all communication using AES-GCM with shared secret T.

7.7 Security Considerations

All authentication security is provided by the secrecy of P during the initial exchange. If the external channel used to communicate P is actively intercepted and an intermediary server is malicious, a MITM attack can be conducted by an adversary. However, this must be an active attack as disclosure of P after the true B value is known by the client renders this MITM attack impossible.

A malicious server or client could attempt to brute force P. However, every attempt requires interaction with the Host. After a few failed attempts, the Host should generate a new P value. If failed attempts continue, the Host should stop accepting connections all together and report an issue to the user. To prevent DOS (denial of service) attacks and/or a malicious client forcing the Host to regenerate P, the Server should add rate limiting and other DOS protection for Clients.

In order to ensure perfect forward secrecy, new key pairs $(PK_{H/C}/pk_{H/C})$ should be generated for each session.