

Screen View Protocol

Josh Brown

November 25, 2021

1 Abstract

Screen View is an end to end encrypted remote screen viewing and controlling software. This document describes the protocols necessary to make it function.

Contents

1	Abstract	2
2	Definitions	4
3	Remote Visual Display (RVD) Protocol	5
3.1	Definitions	5
3.2	Handshake	5
3.2.1	ProtocolVersion - TCP	5
3.2.2	Initialization	5
3.3	Control messages	5
3.3.1	DisplayChange - TCP	6
3.3.2	DisplayChangeReceived - TCP	7
3.3.3	MouseLocation - TCP/UDP	7
3.4	Input	7
3.4.1	MouseInput - TCP/UDP	7
3.4.2	KeyInput - TCP/UDP	8
3.5	Clipboard	8
3.5.1	ClipboardTypeRequest - TCP	8
3.5.2	ClipboardTypeResponse - TCP	8
3.5.3	CopyRequest - TCP	8
3.5.4	CopyResponse - TCP	9
3.5.5	A note on Pasting	9
3.6	FrameData - UDP	9
4	Weak Pre Shared Key, Key Authentication (WPSKKA) Protocol	10
4.1	Introduction	10
4.2	Existing Practice	10
4.3	Goal	11
4.4	Requirements	11
4.5	Other Protocol/Algorithms Definitions	11
4.6	Protocol	11
4.7	Security Considerations	13

2 Definitions

The following definitions are used globally throughout the document:

- Host - The user that wants to share their screen to the Client
- Client - The user that wants to view and control the Host's screen
- Server - The intermediary server used for routing and proxying data between

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in [RFC 2119](#).

3 Remote Visual Display (RVD) Protocol

The RVD protocol is used to communicate mouse input, keyboard input, frame data, and clipboard data between the Host and the Client.

All messages can occur over either TCP or UDP but it is strongly RECOMMENDED that the noted transport protocol is used.

With the exception of the *Handshake* messages. All RVD messages' first byte contain a number to indicate the message type.

3.1 Definitions

- Display - A rectangular visual region that is shared by a Host to a Client. May or may not be *Controllable*.
- Controllable - A *Display* that accepts keyboard and mouse input.

3.2 Handshake

3.2.1 ProtocolVersion - TCP

Handshaking begins by the Host sending the client a *ProtocolVersion* message. This lets the Client know the version supported by the Host.

The *ProtocolVersion* message consists of 12 bytes interpreted as a string of ASCII characters in the format "RVD xxx.yyy" where xxx and yyy are the major and minor version numbers, padded with zeros.

Host → Client
version (11 bytes) - "RVD 001.000"

The client replies back either 0 to indicate the version is not acceptable and that the handshake has failed or 1 if the version is acceptable to the client and the handshake has succeeded. If 0 is sent, all communication MUST cease and an error SHOULD be displayed to user.

Client → Host
ok (1 byte) - 00000000 or 00000001

3.2.2 Initialization

Once the handshake has succeeded the Host responds with a *DisplayChange* message.

3.3 Control messages

Control messages are messages that instruct client about changes regarding the Host.

3.3.1 DisplayChange - TCP

A *DisplayChange* message informs the client about the available *Displays*. RVD supports up to 255 displays.

Host → Client	
type (1 byte)	- 1
clipboard-readable (1 byte)	- 0 or 1
number-of-displays (1 byte)	- 1-255
displays-information (variable bytes)	
DisplayInformation <i>number-of-displays</i> times	

Each *Display* has an associated *DisplayInformation*. *displays-information* contains *number-of-displays* *DisplayInformation*'s. A *DisplayInformation* is defined below:

display-id (1 byte)
width (2 bytes) - number of pixels of the width of this display
height (2 bytes) - number of pixels of the width of this display
cell-width (2 bytes) - number of pixels of the width of a cell in the grid
cell-height (2 bytes) - number of pixels of the height of a cell in the grid
access (1 byte) - defined below
name-length (1 byte) - length of the name
name (<i>name-length</i> bytes) - display name (UTF-8)

Restrictions:

- *cell-width* MUST be less than *width*. *cell-height* must be less than *height*.
- The *access* byte defines what type of access is available for the display. The bits of the *access* byte are described below in Big Endian.

Bit	Description
0	Flush
1	<i>Controllable</i>
2	Reserved for future use
3	
4	
5	
6	
7	

If the *Controllable* bit is 1 and the *clipboard-readable* byte is set to 1, then the clipboard is writable. The *Controllable* bit SHOULD be consistent throughout all displays.

The *Flush* bit indicates whether this display has changed, specifically if this *display-id* refers to a different *Display* than the same *display-id* did in the

previous *DisplayChange* message. In initialization, this MUST always be 1 (as there is no previous *DisplayChange*). If the display hasn't changed (0) then the frame data may be maintained. If *Flush* is 0, *width*, *height* MUST remain the same as the previous *DisplayChange* specified for the *display-id*.

3.3.2 DisplayChangeReceived - TCP

The *DisplayChangeReceived* message is sent in reply after receiving a *DisplayChange* message. It indicates to the Host they may start sending *FrameData* referencing the new *DisplayInformation* in the most recent *DisplayChange*.

Client → Host	
type (1 byte) - 2	

3.3.3 MouseLocation - TCP/UDP

The *MouseLocation* message send information about where the mouse is currently on the screen. The Host sends this information periodically throughout the session. The Host SHOULD send a *MouseLocation* update when mouse input is received from the Host's system or in reply when it receives a *MouseInput*.

Host → Client	
type (1 byte) - 3	
display-id (1 byte) - 0-255	
x-location (2 bytes) - x coordinate of the mouse	
y-location (2 bytes) - y coordinate of the mouse	

3.4 Input

Input messages (including *MouseLocation*) may be sent over TCP or UDP. TCP is preferred in most situations. However, in situations where speed is prioritized over the guarantees TCP provides (such as gaming), UDP can be used.

3.4.1 MouseInput - TCP/UDP

Client → Host	
type (1 byte) - 4	
display-id (1 byte) - 0-255	
x-position (2 bytes) - x coordinate of the mouse	
y-position (2 bytes) - y coordinate of the mouse	
button-mask (1 byte) - described below	

Indicates either pointer movement or a pointer button press or release. The pointer is now at (x-position, y-position), and the current state of buttons 1 to 8 are represented by bits 0 to 7 of button-mask respectively, 0 meaning up, 1 meaning down (pressed).

On a conventional mouse, buttons 1, 2 and 3 correspond to the left, middle and right buttons on the mouse. On a wheel mouse, each step of the wheel is represented by a press and release of a certain button. Button 4 means up, button 5 means down, button 6 means left and button 7 means right.

3.4.2 KeyInput - TCP/UDP

The *KeyInput* event sends key presses or releases.

Client → Host	
type (1 byte) - 5	
down-flag (1 byte) - 0 or 1 to indicate whether the key is now pressed or released	
key (4 bytes) - "keysym"	

Details can be found at the [RFB Spec](#)

3.5 Clipboard

3.5.1 ClipboardTypeRequest - TCP

Used to request clipboard types the Host supports.

Client → Host
type (1 byte) - 6

3.5.2 ClipboardTypeResponse - TCP

Response to the *ClipboardTypeRequest*

Host → Client	
type (1 byte) - 7	
number-of-clipboard-types (1 byte) - 0-255	
data (variable bytes) - described below	

number-of-clipboard-types is always 0 if *clipboard-readable* is 0.

data contains *number-of-clipboard-types* *ClipboardTypes*. *ClipboardType* is defined below.

type-length (1 byte) - 1-255
type-name(<i>type-length</i> bytes) - type name in ASCII

3.5.3 CopyRequest - TCP

This is a request for a keyboard contents. It can be made by either the Client or the Host.

Client ↔ Host	
type (1 byte) - 8	
type-length (1 byte) - 0-255	
type-name(<i>type-length</i> bytes) - type name in ASCII	

3.5.4 CopyResponse - TCP

CopyResponse message is a response to a *CopyRequest*.

Client ↔ Host
type (1 byte) - 9
accepted (1 byte) - 0 or 1
Below only if <i>accepted</i> is 1
type-length (1 byte) - 0-255
type-name(<i>type-length</i> bytes) - type name in ASCII
content-length (4 bytes) - the length of the content (maximum 2^{24} bytes or 16MB)
data (content-length bytes)

accepted indicates whether the *CopyRequest* was accepted. If 0, the rest of the message MUST not exist. If *clipboard-readable* is 0, *accepted* is always 0. A Client or Host may send this message without a request. If a *CopyResponse* is unsolicited, then *accepted* MUST be 1.

data is zlib compressed.

3.5.5 A note on Pasting

There is a no paste message. To paste data an unsolicited *CopyResponse* may be sent and then the keyboard shortcut (ctrl+v or cmd+v) should be sent via the *KeyboardMessage*

3.6 FrameData - UDP

The *FrameData* message contains an update of a particular cell on a particular *Display*.

type (1 byte) - 10
sequence-number (4 bytes)
cell-number (2 bytes)
size (2 bytes)
data (<i>size</i> bytes)

sequence-number is an incrementing 32-bit counter for each *FrameData* sent
data contains jpeg pixel data of the updated cell.

4 Weak Pre Shared Key, Key Authentication (WPSKKA) Protocol

4.1 Introduction

This protocol is used to establish end to end encryption between the Host and the Client.

In recent years end-to-end encryption has risen in popularity due to privacy and security concerns. However, many implementations of end-to-end encryption rely on a third party and/or are susceptible to man-in-the-middle attacks making them inadequate.

Screen sharing applications are used for a multitude of different purposes. One common use case is an IT professional assisting somebody by remotely viewing and controlling their computer. Sensitive data could be visible on the user's screen. Therefore, end-to-end encryption is preferred.

Additionally, IT professionals will often be communicating with the user via the telephone. This provides a bi-directional external channel to transfer information. However, people cannot and will not transfer large amounts or complicated data reliably. Simply communicating letters can be confusing. "B", "C", "D", "E", "G" all sound similar and can be confused. Therefore, communicating strictly numbers is ideal.

This creates an issue. Short, purely numeric keys provide extremely low entropy. A 10 digit numerical code only provides about 33-bits of entropy. In fact, in order to get the ideal 128-bits of entropy a 39 digit key would need to be used. Users will not want to relay 39 digits over the phone.

One common solution when using a weak password is using a KDF to perform key-stretching. However, a 10 digit numerical code has such a small key space that it is relatively easy to brute force. Government agencies, such as the NSA, or large companies, such as Microsoft or Google could easily brute-force even very slow KDFs such as argon2 or PBKDF2.

4.2 Existing Practice

Based on [TeamViewer's security statement](#), TeamViewer is end-to-end encrypted. However, a rogue or malicious TeamViewer intermediary server could easily provide a different set of keys to each party and intercept all communication. The parties must trust this third party in order to achieve proper end-to-end encryption. TeamViewer also does not give the ability for users to check the fingerprint of the other party's public key.

Other applications such as Zoom and Signal (and the Signal protocol itself) do give users the ability to verify some sort of fingerprint of the public key ([Zoom](#) and [Signal](#)). However, most users don't bother confirming the numbers. Again trusting the third party server.

Solutions such as TLS do provide strong end-to-end encryption but rely on a third party Certificate Authority to sign public keys. This would not be possible in the aforementioned use case.

4.3 Goal

Elliptic Curve Diffie-Hellman provides a sufficiently secure means of arriving at a shared secret. The goal of this protocol is to authenticate Elliptic Curve public keys using a weak (3 bytes length, 24 bits of entropy) pre-shared key communicated via an external channel. The security of the external channel is out of scope for this protocol and will be assumed to be secure (see Security Considerations section).

4.4 Requirements

Communication between the parties should be minimized.

The protocol should be secure in the case of a malicious Client and/or malicious Server. This means the protocol should not be susceptible to man-in-the-middle (MITM) attacks, and the Host should be able to self authenticate the Client without trusting the Server.

4.5 Other Protocol/Algorithms Definitions

The Secure Remote Password (SRP) protocol is defined in [RFC2945](#) and [RFC5054](#).

Elliptic Curve Diffie-Hellman (ECDH) key exchange is described in [RFC6090](#).

4.6 Protocol

This protocol occurs after a connection is established between the Host and Client using the Server.

The SRP group is the 2048-bit group from RFC5054:

The hexadecimal value for the prime is:

```

AC6BDB41 324A9A9B F166DE5E 1389582F AF72B665 1987EE07 FC319294
3DB56050 A37329CB B4A099ED 8193E075 7767A13D D52312AB 4B03310D
CD7F48A9 DA04FD50 E8083969 EDB767B0 CF609517 9A163AB3 661A05FB
D5FAAAE8 2918A996 2F0B93B8 55F97993 EC975EEA A80D740A DBF4FF74
7359D041 D5C33EA7 1D281E44 6B14773B CA97B43A 23FB8016 76BD207A
436C6481 F1D2B907 8717461A 5B9D32E6 88F87748 544523B5 24B0D57D
5EA77A27 75D2ECFA 032CFBDB F52FB378 61602790 04E57AE6 AF874E73
03CE5329 9CCC041C 7BC308D8 2A5698F3 A8D0C382 71AE35F8 E9DBFBB6
94B5C803 D89F7AE4 35DE236D 525F5475 9B65E372 FCD68EF2 0FA7111F
9E4AFF73

```

The generator is: 2.

The Host generates the following:

- PK_H/pk_H - Host ephemeral elliptic curve Public/Private key using the `secp521r1` (P-521) curve
- I - 128 bit cryptographic secure random number, used as the identity or username in SRP
- S - SRP salt
- P - 3 byte random cryptographic secure random number, used as the password in SRP
- V - SRP verifier
- b - SRP random private value
- B - SRP public value
- k - SRP K value

The Host sends S, I, and B to the Client.

The Client generates:

- PK_C/pk_C - Client ephemeral elliptic curve Public/Private key using the `secp521r1` (P-521) curve
- a - SRP random private value
- A - SRP public value
- u - SRP u value
- k - SRP k value

- x - SRP x value (hashed P value communicated externally to the client)
- L - SRP session key

The Client sends the Host A , PK_C and $\text{HMAC}(PK_C, L)$.

The Host derives:

- u - SRP U value
- L - SRP session key

The Host authenticates PK_C HMAC using L . If authentication is successful, the Host performs DHKE to derive, T the shared secret.

The Host sends the client PK_H and $\text{HMAC}(PK_H, L)$.

The Client authenticates PK_H HMAC using L . If authentication is successful, the Client performs DHKE to derive, T the shared secret.

Both the Client and the Host encrypt all communication using AES-GCM with shared secret T .

4.7 Security Considerations

All authentication security is provided by the secrecy of P during the initial exchange. If the external channel used to communicate P is actively intercepted and an intermediary server is malicious, a MITM attack can be conducted by an adversary. However, this must be an active attack as disclosure of P after the true B value is known by the client renders this MITM attack impossible.

A malicious server or client could attempt to brute force P . However, every attempt requires interaction with the Host. After a few failed attempts, the Host should generate a new P value. If failed attempts continue, the Host should stop accepting connections all together and report an issue to the user. To prevent DOS (denial of service) attacks and/or a malicious client forcing the Host to regenerate P , the Server should add rate limiting and other DOS protection for Clients.

In order to ensure perfect forward secrecy, new key pairs ($PK_{H/C}/pk_{H/C}$) should be generated for each session.