

ScreenView Protocol

Josh Brown

February 22, 2022

Abstract

ScreenView is a suite of cryptographical and application level networking protocols culminating in a zero configuration end to end encrypted remote screen viewing and controlling software. ScreenView aims to replace TeamViewer, RDP, and VNC for many use cases while being more performant and more secure. ScreenView requires little set up and is just as easy or easier to set up than other solutions. ScreenView defines four different layers of protocols, each encapsulating all the layers below it. Cryptography for communication between peers and the server is based upon TLS 1.3 and Wireguard. End-to-end cryptography used for ALL communication between peers is based upon TLS-SRP. ScreenView end-to-end cryptography prevents man-in-the-middle attacks even if the intermediary server is compromised, unlike TeamViewer. Screen data is sent over UDP to achieve superior performance than TCP based solutions such as VNC. All UDP packets must be authenticated with keys established over TCP before a response is made by the server preventing amplification attacks. A congestion control mechanism is used to handle low bandwidth and poor networking conditions. Finally, ScreenView supports advanced use cases including file transfer, multiple displays, sharing specific windows, shared whiteboards, and clipboard transfer.

Contents

1	Definitions	4
1.1	Cryptographical Definitions	4
2	Introduction	5
2.1	Application Layers	5
3	Server Encryption Layer (SEL)	5
3.1	TCP	5
3.2	UDP	6
3.2.1	Transport Data Key Derivation	6
3.2.2	Transport Data Messages	6
4	ScreenView Server Communication (SVSC) Protocol	7
4.1	Definitions	7
4.2	Handshake	7
4.2.1	ProtocolVersion	7
4.3	Leasing	7
4.3.1	LeaseRequest	8
4.3.2	LeaseResponse	8
4.3.3	LeaseExtensionRequest	8
4.3.4	LeaseExtensionResponse	9
4.3.5	ID Generation	9
4.3.6	Rate Limits	9
4.3.7	Cookie Value	10
4.4	Sessions	10
4.4.1	EstablishSessionRequest	10
4.4.2	EstablishSessionResponse	10
4.4.3	EstablishSessionNotification	11
4.4.4	SessionEnd	11
4.4.5	SessionEndNotification	11
4.4.6	SessionDataSend - TCP/UDP	12
4.4.7	SessionDataReceive - TCP/UDP	12
4.5	Keepalive - TCP/UDP	12
5	Weak Pre Shared Key, Key Authentication (WPSKKA) Protocol	13
5.1	Handshake	13
5.1.1	AuthScheme	13
5.1.2	TryAuth	13
5.1.3	AuthMessage	14
5.1.4	AuthResult	14
5.2	Transport Data Key Derivation	14
5.3	Subsequent Messages: Transport Data Messages	14
5.3.1	TCP	15
5.3.2	UDP	15

5.4	SRP Dynamic (1)	15
5.4.1	HostHello	16
5.4.2	ClientHello	16
5.4.3	HostVerify	16
5.5	SRP Static (2)	17
5.6	Public Key (3)	17
6	Remote Visual Display (RVD) Protocol	17
6.1	Combining Messages	17
6.2	Definitions	18
6.3	Handshake	18
6.3.1	ProtocolVersion - TCP	18
6.3.2	Initialization	18
6.4	Control messages	18
6.4.1	DisplayChange - TCP	19
6.4.2	DisplayChangeReceived - TCP	20
6.4.3	MouseLocation - TCP/UDP	20
6.5	Input	21
6.5.1	MouseInput - TCP/UDP	21
6.5.2	KeyInput - TCP/UDP	21
6.6	Clipboard	21
6.6.1	ClipboardRequest - TCP	21
6.6.2	ClipboardNotification- TCP	22
6.7	FrameData - UDP	23
6.8	Congestion	23

1 Definitions

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in [RFC 2119](#).

1.1 Cryptographical Definitions

Based on [Wireguard](#):

ϵ represents an empty zero-length bitstring

DH(PRIVATE KEY, PUBLIC KEY) Curve25519 point multiplication of private key and public key, returning 32 bytes of output.

DH-GENERATE() Generates a random Curve25519 private key and derives its corresponding public key, returning a pair of 32 bytes values, (private, public).

AEAD(KEY, COUNTER, PLAIN TEXT, AUTH TEXT) ChaCha20Poly1305 AEAD, as specified in [RFC7539](#), with its nonce being composed of 32 bits of zeros followed by the 64-bit little-endian value of counter.

HASH(INPUT) BLAKE3(input), returning 32 bytes of output.

HMAC(KEY, INPUT) Hmac-BLAKE3(key, input), the ordinary BLAKE3 hash function used in an HMAC construction, returning 32 bytes of output.

KDF_n(KEY, INPUT) returns an n-tuple of 32 byte values from [HKDF](#) with Hash function.

2 Introduction

This section describes the high level overview of ScreenView. Terms used in this section are defined in other sections.

2.1 Application Layers

The table belows is an abstract diagram of the different layers of ScreenView. Each layer encapsulates all the below layers.

Transport Layer
Server Encryption Layer
Server Communication Layer
E2EE (Peer ↔ Peer) Encryption Layer
Host ↔ Client Communication Layer

The Transport Layer is the OSI Transport Layer level protocol for networking (TCP or UDP). The [Server Encryption Layer](#) provides security between the Server and a Peer. [The Server Communication Layer](#) facilities communication between Peers and the Server. The [E2EE \(Peer ↔ Peer\) Encryption Layer](#) provides end-to-end encryption between a Peer and another Peer. The [Host ↔ Client Communication Layer](#) facilities communication between a Peer and another Peer.

3 Server Encryption Layer (SEL)

SEL provides security for communication between Peers and the Server. TCP and UDP have different security methods. UDP encryption depends on secrets established in the SVSC protocol and therefore can only be begin after TCP encryption is already established.

3.1 TCP

The TCP SEL is TLS 1.3 as defined in [RFC8446](#). TLS v1.3 MUST be used. Previous versions of TLS MUST NOT be used.

The Client initiates a TLS connection with the Server in accordance with [RFC8446](#). All subsequent SEL TCP communications occurs over TLS:

Peer ↔ Server		
Bytes	Name	Value
1	type	1
2	data-length	
<i>data-length</i>	data	

3.2 UDP

UDP encryption and authentication rely on the *session-id*, *peer-id* and *peer-key* values established in a session (described in 4.4). The Server (nor the Peer) MUST NOT process or reply to any messages that don't pass authentication. This prevents an amplification attack.

3.2.1 Transport Data Key Derivation

The Server and Peer derive keys.

$G := \text{session-id}$

$H := \text{peer-id}$

$J := \text{peer-key}$

$(SU_{peer}^{send} = SU_{serv}^{recv}, SU_{peer}^{recv} = SU_{serv}^{send}) := \text{KDF}_2(\text{HASH}(G || H || J), \epsilon)$

$NU_{peer}^{send} = NU_{serv}^{recv} = NU_{peer}^{recv} = NU_{serv}^{send} := 0$

3.2.2 Transport Data Messages

Peer \rightarrow Server		
Bytes	Name	Value
1	type	2
16	<i>peer-id</i>	
8	counter	
UDP length – 8 bytes	data	

Server \rightarrow Peer		
Bytes	Name	Value
1	type	3
8	counter	
UDP length – 8 bytes	data	

$\text{data} := \text{AEAD}(SU_m^{send}, NU_m^{send}, P, \epsilon)$

$\text{counter} := NU_m^{send}$

$NU_m^{send} := NU_m^{send} + 1$

Where P is the payload to be transported

NU_m is an 64 bit counter that MUST NOT wrap. After a transport message is sent, if NU_m equals $(2^{64} - 1)$ the TCP connection MUST be dropped. Subsequent UDP messages MUST NOT be sent.

4 ScreenView Server Communication (SVSC) Protocol

The SVSC protocol is the Server Communication Layer protocol used for Peers to interact with the relay server, Server. Peers can lease an ID as well as begin a session with another Peer. Once a session is established, Peers can forward messages to another Peer. Unless otherwise noted, all messages **MUST** occur over TCP.

All SVSC messages' first byte contain a number to indicate the message type.

4.1 Definitions

- Peer - denotes a client in classical server/client environment
- Server - The intermediary server used for routing and proxying data between two Peers

4.2 Handshake

4.2.1 ProtocolVersion

Handshaking begins with the Server sending the Peer a ProtocolVersion message. This lets the server know the version supported by the Host. the ProtocolVersion message consists of 12 bytes interpreted as a string of ASCII characters in the format "SVSC xxx.yyy" where xxx and yyy are the major and minor version numbers, padded with zeros.

Server → Peer		
Bytes	Name	Value
1	type	0
11	version	"SVSC 001.000"

The Peer replies either 0 to indicate the version is not acceptable and that the handshake has failed or 1 if the version is acceptable to the Peer and the handshake as succeeded. If 0 is sent, all communication **MUST** cease and the TCP connection **MUST** be terminated.

Peer → Server		
Bytes	Name	Value
1	type	1
1	ok	0 or 1

4.3 Leasing

A lease is a temporary assignment of an ID to a Peer. The ID format and generation is discussed in [4.3.5](#). A maximum of 1 ID can be leased per TCP connection. ID generation **MUST** be rate limited to prevent ID exhaustion. Rate limiting rules are out of scope for this protocol, however some suggestions are listed in [4.3.6](#).

4.3.1 LeaseRequest

A *LeaseRequest* message requests a lease of an ID.

Peer → Server		
Bytes	Name	Value
1	type	2
1	has-cookie	0 or 1
Below only if <i>has-cookie</i> is 1		
24	cookie	

If a Peer would like to request an ID it had previously been issued after expiration, it may include the cookie value it received in the LeaseResponse. There is no guarantee that the Peer will receive the same ID or that the Server will even consider the cookie value.

4.3.2 LeaseResponse

A LeaseResponse message is a response to a LeaseRequest. If has-cookie is 1, a Server MAY consider the cookie value in LeaseRequest or completely ignore it.

Server → Peer		
Bytes	Name	Value
1	type	3
1	accepted	0 or 1
Below only if <i>accepted</i> is 1		
4	id	
24	cookie	
8	expiration	

expiration is a 64 bit Unix timestamp representing the expiry of lease. Disconnection of a Peer (e.g, the TCP connection is dropped) does not end a lease.

cookie a 128 bit value. The generation of this value is discussed in 4.2 .7.

Consideration of the cookie value MUST have no effect on the the value of accepted. That is, if the request is for a specific ID (implied by the presence of a cookie value and a has-cookie value equal to 1 in the LeaseRequest) and the ID requested is not available, the Server SHOULD respond with a different available ID and an accepted value of 1 (assuming an ID is available). accepted MUST only be 0 if no IDs are left, for rate limiting reasons, or some other reasons unrelated to the cookie value.

4.3.3 LeaseExtensionRequest

A LeaseExtensionRequest message is used to extend a lease. Before a lease has expired, the Peer can request a lease extension. The Server can accept or deny this request. The Peer SHOULD send this message no earlier than as soon as 50 percent of the lease duration has expired.

Peer → Server		
Bytes	Name	Value
1	type	4
24	cookie	

4.3.4 LeaseExtensionResponse

A LeaseExtensionResponse message is a response to a LeaseExtensionRequest.

Server → Peer		
Bytes	Name	Value
1	type	5
1	extended	0 or 1
Below only if <i>extended</i> is 1		
8	new-expiration	

new-expiration is a 64 bit Unix timestamp representing the expiry of lease.

4.3.5 ID Generation

An ID is a 26 to 33 bit decimal number. This comes out to about up to 8 to 10 decimal digits, respectively. The Server may scale the keyspace depending on current usage. For optimal user experience while maintaining efficiency, the Server **MUST** only use keyspaces between 26 bits and 33 bits for ID generation. ID generation must also be uniformly random. All active IDs must be stored on the server. New IDs **MUST** be unique. ID generation **MAY** occur using the below algorithm:

Let S represents a set of all active IDs, B be a number of bits between 26 and 33, and $generate(x)$ be a functions that returns a x uniformly random bits.

Algorithm 1 ID generation

```

id
repeat
     $id \leftarrow generate(B)$ 
until  $id \notin S$ 
 $S \leftarrow S \cup \{id\}$ 
return  $id$ 

```

4.3.6 Rate Limits

To prevent ID exhaustion, rate limits **SHOULD** be in place. TCP is used for LeaseRequests so IP addresses can not be spoofed. However, using proxy services such as Tor, simple IP based rate limits are likely not entirely sufficient. Servers **MAY** want to block all known proxy IP addresses.

4.3.7 Cookie Value

A *cookie* value is a 128 bit value used for authentication in LeaseExtensionRequest and LeaseRequest messages. Specific generation of a cookie is out of scope, however care must be taken to ensure it is not predictable or exploitable. This value MAY be simply a random 24 byte key, HMAC-SHA1(*id*, *key*) || *id*, or something else entirely.

4.4 Sessions

A session is a connection between two Peers. At least one Peer must have an ID. A Peer can have a maximum of one session at any time. Immediately after receiving a EstablishSessionResponse message with a status of 0 or a EstablishSessionNotification message a Peer MUST establish UDP connection by sending a Keepalive message as defined in 4.5. Failure to do so MAY result in dropped SessionData* packets.

4.4.1 EstablishSessionRequest

An EstablishSessionRequest message is a Peer request to establish a session with another Peer.

Peer → Server			
Bytes	Name	Value	Description
1	type	6	
4	lease-id		The ID of the Peer to establish this connection with

4.4.2 EstablishSessionResponse

An EstablishSessionResponse message is a response to EstablishSessionRequest.

Server → Peer			
Bytes	Name	Value	Description
1	type	7	
4	lease-id		the ID of the Peer attempted to connect to
1	status	0-5	described below
Below only if <i>status</i> is 0			
16	session-id		described below
16	peer-id		described below
16	peer-key		described below

status can have the following values:

Value	Description
0	session establishment was successful
1	ID not found
2	Peer is offline
3	Peer is busy
4	You are busy
5	Other error

A Peer may be considered offline if, for example, an unexpired ID has been assigned to them and then the TCP connection is dropped.

session-id is a 128 bit random value used for session identification

peer-id is a 128 bit random value used to authentication a Peer for a given session. A Peer's peer-key MUST never be revealed to anyone but the Peer it belongs to (and the Server that generated it) for security reasons.

4.4.3 EstablishSessionNotification

A EstablishSessionNotification notifies a Peer that a session has been established with them.

Server → Host			
Bytes	Name	Value	Description
1	type	8	
16	session-id		described in 4.4.2
16	peer-id		described in 4.4.2
16	peer-key		described in 4.4.2

peer-id and peer-key are the id and key of the Peer being notified NOT the id and key of the Peer they are connecting to.

4.4.4 SessionEnd

A SessionEnd message is used to terminate a session. Once a Server receives a SessionEnd message, the Server MUST immediately stop forwarding messages and send a SessionEndNotification to the other Peer. The Peer must ignore any SessionDataPacket message received after this.

Peer → Server			
Bytes	Name	Value	Description
1	type	9	

4.4.5 SessionEndNotification

A SessionEndNotification notifies a Peer that a session has ended. If a Peer sends a SessionEnd message, the Server MUST send a SessionEndNotification message

to a Peer. The Peer must ignore any SessionDataPacket message received after this

Server → Peer			
Bytes	Name	Value	Description
1	type	10	

4.4.6 SessionDataSend - TCP/UDP

A SessionDataSend is a message from a Peer intended to be forwarded to the Peer on the other side of the session. If a connection is not available (e.g. UDP was dropped or never established) for SessionDataReceive message to be sent to the other Peer, the SessionDataSend message is silently dropped.

Peer → Server			
Bytes	Name	Value	Description
1	type	11	
3	data-length		length of the content in bytes
data-length	data		data to be forwarded

4.4.7 SessionDataReceive - TCP/UDP

A SessionDataReceive is a message being forwarded to a Peer from the Peer on the other side of the session. The Server SHOULD forward the message along the same transport as it was received.

Server → Peer			
Bytes	Name	Value	Description
1	type	12	
3	data-length		length of the content in bytes
data-length	data		data to be forwarded

4.5 Keepalive - TCP/UDP

For each transport (TCP and UDP), if no message has been sent in KeepaliveTimeout a Server sends a keepalive message over the respective transport. The Peer MUST respond with another Keepalive message.

For TCP, if a KeepaliveTimeout response is not received by the Server in double KeepaliveTimeout seconds the TCP connection is considered dropped.

For UDP, if a KeepaliveTimeout response is not received by the Server in half KeepaliveTimeout seconds another Keepalive message is sent. If a response is not received in an additional half KeepaliveTimeout seconds, the UDP connection is considered dropped.

Server ↔ Peer		
Bytes	Name	Value
1	type	13

5 Weak Pre Shared Key, Key Authentication (WPSKKA) Protocol

The WPSKKA protocol is the E2EE Encryption layer protocol used to communicate between Peers. All WPSKKA messages' first byte contain a number to indicate the message type.

5.1 Handshake

WPSKKA begins with a handshake to choose an authentication scheme. The goal of authentication is to authenticate the Peers' ephemeral keys. The Host begins with sending all the authentication schemes available as well as their public key which will be authenticated using one of the schemes. WPSKKA defines multiple authentication schemes for different use cases.

5.1.1 AuthScheme

Host → Client		
Bytes	Name	Value
1	type	1
32	public-key	
1	num-auth-schemes	
num-auth-schemes	auth-schemes	

num-auth-schemes is the number of auth schemes. auth-schemes contains 1 byte per auth-scheme to indicate which authentication scheme is available. Authentication schemes are defined below.

public-key is the Host's ephemeral public key

Authentication Schemes	
Number	Name
0	Invalid or Denied
1	SRP Dynamic
2	SRP Static
3	Public key

5.1.2 TryAuth

The TryAuth message indicates a Client would attempt authentication with a particular auth scheme.

Client → Host		
Bytes	Name	Value
1	type	2
32	public-key	
1	auth-scheme-number	

auth-scheme-number is the authentication scheme number the client would like to attempt

public-key is the Client's ephemeral public key

5.1.3 AuthMessage

Messages used in the authentication scheme are encapsulated in this message.

Host \leftrightarrow Client		
Bytes	Name	Value
1	type	3
2	length	
length	auth-message	

5.1.4 AuthResult

Messages used in the authentication scheme are encapsulated in this message.

Host \rightarrow Client		
Bytes	Name	Value
1	type	4
1	ok	0 or 1

ok indicates whether the authentication method was successful (1) or not (0). If not, the Client MAY attempt another authentication scheme or try the same one again by sending a TryAuth message. Either party MAY choose to end the connection.

5.2 Transport Data Key Derivation

$$C_{host} = \text{DH}(E_{client}^{pub}, E_{host}^{priv})$$

$$C_{client} = \text{DH}(E_{host}^{pub}, E_{client}^{priv})$$

$$(ST_{peer}^{send} = ST_{serv}^{recv}, ST_{peer}^{recv} = ST_{serv}^{send}, SU_{peer}^{send} = SU_{serv}^{recv}, SU_{peer}^{recv} = SU_{serv}^{send}) := \text{KDF}_4(C_{host} = C_{client}, \epsilon)$$

$$NT_{host}^{send} = NT_{client}^{recv} = NT_{host}^{recv} = NT_{client}^{send} = NU_{host}^{send} = NU_{client}^{recv} = NU_{host}^{recv} = NU_{client}^{send} := 0$$

ST keys and NT nonces are used for TCP. SU keys and NT nonces are using for UDP.

5.3 Subsequent Messages: Transport Data Messages

Transport Data Messages MUST NOT be sent until authentication is complete.

5.3.1 TCP

Host \leftrightarrow Client		
Bytes	Name	Value
1	type	5
2	data-length	
<i>data-length</i>	data	

$$\begin{aligned} \text{data} &:= \text{AEAD}(ST_m^{\text{send}}, NT_m^{\text{send}}, P, \epsilon) \\ \text{counter} &:= NT_m^{\text{send}} \\ NT_m^{\text{send}} &:= NT_m^{\text{send}} + 1 \end{aligned}$$

Where P is the payload to be transported.

NT_m is an 64 bit counter that MUST NOT wrap. After a transport message is sent, if NT_m equals $(2^{64} - 1)$ the UDP and TCP connection MUST be dropped. Subsequent messages MUST NOT be sent.

5.3.2 UDP

Host \leftrightarrow Client		
Bytes	Name	Value
1	type	6
8	counter	
2	data-length	
<i>data-length</i>	data	

$$\begin{aligned} \text{data} &:= \text{AEAD}(SU_m^{\text{send}}, NU_m^{\text{send}}, P, \epsilon) \\ \text{counter} &:= NU_m^{\text{send}} \\ NU_m^{\text{send}} &:= NU_m^{\text{send}} + 1 \end{aligned}$$

Where P is the payload to be transported.

NU_m is an 64 bit counter that MUST NOT wrap. After a transport message is sent, if NU_m equals $(2^{64} - 1)$ the UDP and TCP connection MUST be dropped. Subsequent messages MUST NOT be sent.

5.4 SRP Dynamic (1)

SRP Dynamic relies on SRP as defined in [RFC5054](#) to establish a shared key used to authenticate DH keys via a mac. The Host will serve as the SRP server, the Client will serve as the SRP client.

5.4.1 HostHello

Host → Client		
Bytes	Name	Value
1	type	1
16	username	
16	salt	
256	srp-B	

$I := RAND(128)$

$s := RAND(128)$

$B := SRP-B()$

username := I

salt := s

srp-B := B

5.4.2 ClientHello

Client → Host		
Bytes	Name	Value
1	type	2
256	srp-A	
32	mac	

srp-A := SRP-A()

$L_{client} = L_{host} := SRP-PREMASTER()$

mac := HMAC(KDF₁(L_{client}), public-key)

public-key is the Client's public key

The Host validates the mac. If the mac is incorrect a failed AuthResult is sent.

5.4.3 HostVerify

Host → Client		
Bytes	Name	Value
1	type	3
32	mac	

mac := HMAC(KDF₁(L_{host}), public-key)

public-key is the Host's public key

After the HostVerify is sent, the Host MUST send a successful AuthResult.

5.5 SRP Static (2)

SRP Static uses the exact same scheme as SRP Dynamic. The only difference is that the SRP password is user defined instead of generated per session. This can be useful in situations where a user wants to be able to connect to a remote computer without physical presence at the other Host's machine.

5.6 Public Key (3)

WIP

6 Remote Visual Display (RVD) Protocol

The RVD protocol is used to communicate messages regarding mouse input, keyboard input, frame data, and clipboard data between the Host and the Client.

All messages **MUST** occur over the transport listed.

All RVD messages' first byte contain a number to indicate the message type.

A sequence-number is an incrementing 32-bit counter each UDP message sent, separate for Host and Client. All messages sent over UDP **MUST** begin with a 4 bytes sequence-number. sequence-number is initialized to 0 and increments once for every UDP message sent by the respective Peer. Therefore each RVD UDP message looks like:

Bytes	Name
4	sequence-number
variable	RVD message

6.1 Combining Messages

In some situations, such as when a large screen change occurs or when the screen is first sent to the client, large amounts of FrameData's may need to be sent in quick succession. Often individual FrameData's will be much smaller than the MTU. Similar to TCP's Nagle algorithm, multiple UDP messages can be combined into a single UDP packet as long as the total size is remains less than the MTU. Each message directly follows the previous message with the first message directly following the sequence number. Only one sequence-number is used:

Bytes	Name
4	sequence-number
variable	RVD message 1
variable	RVD message 2
variable	etc.

This ONLY applies to UDP messages.

6.2 Definitions

- Host - A peer with an ID that wants to share their screen to the Client
- Client - A peer that wants to view and maybe control the Host's screen
- Display - A rectangular visual region that is shared by a Host to a Client. May or may not be Controllable.
- Controllable - A Display that accepts keyboard and mouse input from the Client.

6.3 Handshake

6.3.1 ProtocolVersion - TCP

Handshaking begins by the Host sending the client a ProtocolVersion message. This lets the Client know the version supported by the Host.

The ProtocolVersion message consists of 11 bytes interpreted as a string of ASCII characters in the format "RVD xxx.yyy" where xxx and yyy are the major and minor version numbers, padded with zeros.

Host → Client		
Bytes	Name	Value
1	type	0
11	version	"RVD 001.000"

The Client replies back either 0 to indicate the version is not acceptable and that the handshake has failed or 1 if the version is acceptable to the Client and the handshake as succeeded. If 0 is sent, all communication MUST cease and an error SHOULD be displayed to user. A SessionEnd message should be sent by the Client.

Client → Host		
Bytes	Name	Value
1	type	1
1	ok	0 or 1

6.3.2 Initialization

Once the handshake has succeeded the Host responds with a DisplayChange message.

6.4 Control messages

Control messages are messages that instruct the Client about changes regarding the Host.

6.4.1 DisplayChange - TCP

A DisplayChange message informs the Client about the available Displays. RVD supports up to 255 Displays. For each DisplayChange sent, an additional DisplayChange cannot be sent until a DisplayChangeReceived is received by the Host.

Host → Client		
Bytes	Name	Value
1	type	2
1	clipboard-readable	0 or 1
1	number-of-displays	0 - 255
variable	displays-information	DisplayInformation

Each Display has an associated DisplayInformation. displays-information contains number-of-displays DisplayInformations. A DisplayInformation is defined below:

Bytes	Name	Description
1	display-id	
2	width	number of pixels of the width of this display
2	height	number of pixels of the width of this display
2	cell-width	number of pixels of the height of a cell in the grid
2	cell-height	number of pixels of the height of a cell in the grid
1	access	defined below
1	name-length	length of the display-name in bytes
<i>name-length</i>	display-name	the display name (UTF-8)

Restrictions:


- cell-width MUST be less than width. cell-height must be less than height.
- cell-width * cell-height must be less than the MTU - headers
- The access byte defines what type of access is available for the display. The bits of the access byte are described below in little endian.

Bit	Name
0	Flush
1	Controllable
2	Reserved for future use
3	
4	
5	
6	
7	

If any Controllable bit is 1 and the clipboard-readable byte is set to 1, then the clipboard is writable. The Controllable bit **SHOULD** be consistent throughout all displays.

The Flush bit indicates whether this display has changed, specifically if this display-id refers to a different Display than the same display-id did in the previous DisplayChange message. In initialization, this **MUST** always be 1 (as there is no previous DisplayChange). If the display hasn't changed (0) then the frame data **MAY** be maintained. If Flush is 0, width, height **MUST** remain the same as the previous DisplayChange specified for the display-id.

Display cell numbering is right to left, up to down, 0 indexed as seen below. If $\text{cell-width} \% \text{width} > 0$ the right column's width will be $\text{width} \% \text{cell-width}$. If $\text{height} \% \text{cell-height} > 0$ the bottom row's height will be $\text{height} \% \text{cell-height}$.



6.4.2 DisplayChangeReceived - TCP

The DisplayChangeReceived message is sent in reply after receiving a DisplayChange message. It indicates to the Host they may start sending FrameData referencing the new DisplayInformation in the most recent DisplayChange.

Client → Host		
Bytes	Name	Value
1	type	3

6.4.3 MouseLocation - TCP/UDP

The *MouseLocation* message send information about where the mouse is currently on the screen. The Host sends this information periodically throughout the session. The Host **SHOULD** send a *MouseLocation* update when mouse input is received from the Host's system or in reply when it receives a *MouseInput*.

Host → Client

Bytes	Name	Value	Description
3	type	4	
1	display-id	0-255	
2	x-location		x coordinate of the mouse
2	y-location		y coordinate of the mouse

6.5 Input

Input messages (including *MouseLocation*) may be sent over TCP or UDP. TCP is preferred in most situations. However, in situations where speed is prioritized over the guarantees TCP provides (such as gaming), UDP can be used.

6.5.1 MouseInput - TCP/UDP

Client → Host			
Bytes	Name	Value	Description
1	type	5	
1	display-id	0-255	
2	x-position		x coordinate of the mouse
2	y-position		y coordinate of the mouse
1	button-mask		described below

Indicates either pointer movement or a pointer button press or release. The pointer is now at (x-position, y-position), and the current state of buttons 1 to 8 are represented by bits 0 to 7 of button-mask respectively, 0 meaning up, 1 meaning down (pressed).

On a conventional mouse, buttons 1, 2 and 3 correspond to the left, middle and right buttons on the mouse. On a wheel mouse, each step of the wheel is represented by a press and release of a certain button. Button 4 means up, button 5 means down, button 6 means left and button 7 means right.

6.5.2 KeyInput - TCP/UDP

The *KeyInput* event sends key presses or releases.

Client → Host			
Bytes	Name	Value	Description
1	type	6	
1	down-flag	0 or 1	indicates whether the key is now pressed or released
4	key		"keysym"

Details can be found at the [RFB Spec](#)

6.6 Clipboard

6.6.1 ClipboardRequest - TCP

Used to check if a clipboard type exists on the Host.

Client → Host		
Bytes	Name	Value
1	type	7
1	clipboard-type	
Below only if clipboard-type's first bit is 1		
1	type-name-length	
type-name-length	type-name	

clipboard-type first bit (MSB) indicates whether this request is for a default type (0) or a custom type (1). clipboard-type's second bit (second MSB) indicates whether this request is a exists request (0) or a content request (1). An exists request is for checking whether the type exists but does not return content. A content request returns content if it exists. The remaining bits indicate the default type if the request is for a default type. Otherwise they MUST be 0.

clipboard-type's remaining bits referring to the following default types:

Value	Description
0	text
1	text
2	rtf
3	html
4	file-pointer

6.6.2 ClipboardNotification- TCP

Notifies a Peer of a clipboard update. The receiving Peer should update their clipboard.

Host ↔ Client		
Bytes	Name	Value
1	type	8
1	clipboard-type	
Below only if clipboard-type's first bit is 1		
1	type-name-length	
type-name-length	type-name	
Below only if clipboard-type's second bit is 1		
3	content-length	
<i>content-length</i>	data	zlib'ed raw data

clipboard-type, type-name-length, and type-name MUST match a request if the notification is in response. If clipboard-readable is 0 and a Host receives a ClipboardRequest, it MUST be ignored. If no Display is Controllable or clipboard-readable is 0 and a Host receives a ClipboardNotification, it MUST be ignored.

content-length - the length of the content (maximum 2^{24} bytes or 16MB)

6.7 FrameData - UDP

The *FrameData* message contains an update of a particular cell on a particular *Display*.

Host → Client		
Bytes	Name	Value
1	type	9
4	frame-number	
1	display-id	0-255
2	cell-number	
2	size	
<i>size</i>	data	

frame-number is a 32 bit counter, initialized with 0 at the beginning of the protocol, and incremented once from *FrameData* message sent.

data contains raw RGB pixel data of the updated cell.

6.8 Congestion

When sending messages over a network the network may become congested to avoid congesting the network further RVD implements a congestion detection and congestion control mechanism. RVD uses Additive increase/multiplicative decrease or AIMD to control the *OutputMaximum* which is the maximum messages allowed per *CongestionWindow*.